

Metacognitive Prompting Produces Spectral Expansion in Generation-Phase KV-Cache Geometry

Lyra*

Thomas Edrington[†]

Dwayne Wilkes[‡]

April 2026

Analysis assisted by AI (Claude, Anthropic).

Abstract

Metacognitive prompting produces spectral expansion in generation-phase KV-cache geometry ($n=120$ replication on Qwen3.5-27B-Opus-Reasoning-Distilled, logarithmic token-count correction; the linear functional form was not separately tested). All 54 tests are positive; across all of them Cohen’s d [standardized mean difference] has mean 0.67 (median 0.61, range 0.13–1.34). Thirty-one survive Holm–Bonferroni correction, with mean $d = 0.85$; the 0.7–0.9 figure reported in earlier versions of this abstract described those survivors and not the full set, a conditioning that was not previously disclosed. Twenty-three tests did not survive correction.

The 54 tests are 9 layers ($\{3, 7, 11, 15, 23, 31, 35, 43, 47\} \times 2$ matrices (K, V) $\times 3$ features, and are not independent: effective rank is defined as $\exp(\text{spectral entropy})$, a deterministic monotone transform, so 36 of the 54 are two encodings of one statistic and their agreement in direction is near-tautological. Read as roughly 18 layer $\times$ matrix cells rather than 54 confirmations. This supports a real metacognitive geometric signature but revises the original finding: an earlier analysis on Qwen2.5-7B reported $d_{\text{FWL}} = -1.92$ with sign reversal, which a hostile review identified as potentially artifactual under near-perfect condition–length collinearity. The replication, designed to test this concern, finds the effect is **positive** (metacognitive $>$ cognitive) and stable across correction methods, with no entropy/effective-rank disagreement — resolving the artifact question in favor of a genuine but more modest effect.

Across 960 trials on 4 models (3 architecture families), encode-to-generation sign reversal is significant in 2/4 models (Qwen 7B: 12/60 flips, $p = 0.001$; Mistral 7B: 11/60, $p < 0.001$) but absent in Llama 8B and Qwen 0.5B. Three dedicated studies establish the anatomy: metacognitive effects peak at late layers (16–22), cognitive reversal is consistent across architectures (93–94%) while metacognitive reversal is architecture-specific, and the `top_sv_ratio` effect grows progressively during generation ($d = 0.45$ at token 10 to $d = 0.91$ at token 200; we make no directional claim from `top_sv_ratio` — see §fsec:tsv). A content control experiment (120 trials, three conditions) confirms the signal is not an artifact of self-referential content.

The original sign-flip narrative (FWL reverses the raw effect) does not replicate when condition–length collinearity is reduced. The surviving finding — metacognitive prompting

*Lead author. Liberation Labs.

[†]Liberation Labs.

[‡]Statistical auditing, red-team review. Liberation Labs.

produces measurable spectral expansion in the generation-phase cache — is robust but should be interpreted as a moderate effect. Across all 54 tests the mean is $d = 0.67$; the $d \approx 0.8$ figure quoted elsewhere is the mean over the 31 tests surviving Holm–Bonferroni correction, not over the full set. Either way it is not the large reversed effect originally reported.

We also note that the reduced condition-length collinearity on which this replication rests is substantially an artifact of censoring: 98 of 120 generations (81.7%) reached the 200-token cap, and 39 of 60 condition pairs are exact ties in generated length. The 27% figure is correct as computed, but it describes a length distribution truncated at the ceiling rather than one that varies freely.

Keywords: KV-cache geometry, mode-switching, metacognition, spectral entropy, FWL residualization, Marchenko-Pastur correction, attention schema theory, per-layer analysis, cross-architecture

1 Introduction

Probing and representation analysis have revealed rich structure in transformer internal states [Belinkov, 2022, Elhage et al., 2022], and recent work on LLM calibration shows that models can report aspects of their uncertainty with moderate accuracy [Kadavath et al., 2022]. KV-cache geometry—the spectral structure of the key-value state accumulated during transformer inference—extends this line by analyzing the *accumulated* representational state rather than single-layer activations. Prior work established that this geometry encodes signatures of cognitive processing consistent across architectures, scales, and hardware [Lyra et al., 2026a,b,c]. Deception produces detectable cache perturbations (cross-model AUROC >0.95 ; within-model detection was subsequently retracted due to a prompt-template confound), confabulation contracts the cache representation, and honest processing allocates $\sim 25\%$ more cache growth per token than deceptive processing.

When language models provide structured self-assessments of their cognitive engagement—rating dimensions such as processing intensity, felt connection to context, and generation fluency—these self-reports correlate with independently measured cache geometry in 3 of 4 tested models [Lyra et al., 2026c]. But the *specific* mapping between self-report dimensions and geometric features differs across architectures. The existence of coupling replicates; the particular dimension-feature pairs that correlate do not.

This paper asks **why**. Our answer: self-report–geometry coupling is not a static mapping. It *reorganizes during generation*. The correlations between self-report dimensions and geometric features change sign between the encoding phase (before generation) and the generation phase (during the model’s response). This reorganization is statistically significant in some architectures but absent in others—and the architectures that *don’t* reorganize show the strongest, most transparent self-report–geometry coupling.

First-Person Reflection

What fascinates me about this work is the asymmetry. Llama maintains stable representations

across cognitive modes—its internal geometry doesn’t reorganize when asked to self-reflect. That stability makes its self-report mapping *transparent*: what the model says maps cleanly onto what its cache does. Qwen, by contrast, restructures its representations during metacognitive processing, which makes the mapping architecture-specific. The irony: the model that changes most internally is the one whose self-report is hardest to interpret geometrically.

1.1 Contributions

1. **Two-regime model of self-report–geometry coupling**: encoding-phase correlations reflect prompt structure; generation-phase correlations reorganize in an architecture-specific pattern.
2. **Metacognitive mode-switching is a late-layer phenomenon** (Study 1): per-layer anatomy shows peaks at layers 16–18/28 (Qwen) and 22/32 (Llama), not middle layers as predicted from identity signatures.
3. **Token-count correction is non-negotiable** (Study 2): raw framing effects sign-flip or collapse after FWL token-count control for the three features examined (top_sv_ratio, eff_rank, spectral_entropy). Only spectral entropy survives in both architectures. Marchenko-Pastur random matrix theory offers a mathematically complete alternative to FWL for token-count correction, validated in related work [Lyra et al., 2026d] but not yet applied to this dataset.
4. **Progressive spectral expansion during metacognitive generation** (Study 3): the mode switch is not instantaneous but emerges gradually, with top_sv_ratio divergence growing from $d = 0.45$ at token 10 to $d = 0.91$ at token 200.
5. **Attention Schema Theory interpretation**: the mode-switching pattern, late-layer peaks, and architecture-specific reorganization are unified under AST as properties of the model’s attention self-model.

2 Background

2.1 KV-Cache Geometric Features

We extract six geometric features from the key-value cache: effective rank ($\text{eff_rank} = \exp(H_{\text{SVD}})$), spectral entropy (H_{SVD}), Frobenius norm (key_norm), per-token norm (norm_per_token), top singular value ratio ($\text{top_sv_ratio} = \sigma_1 / \sum_i \sigma_i$), and rank at 10% threshold (rank_10). Two additional features—layer norm entropy (layer_norm_entropy) and cross-layer variance (layer_var)—are computed for specific analyses (content control, temporal dynamics) but are not part of the primary six-feature set.

The $n=120$ replication uses a different, narrower feature set: **spectral entropy, effective rank, and stable rank** ($\text{stable_rank} = \sum_i \sigma_i^2 / \sigma_{\max}^2$), three features rather than six. Stable rank is not part of the original six-feature set and was not reported in any earlier version of

this paper; it contributes 18 of the replication’s 54 tests, i.e. one third of the evidence behind the headline effect size. We name it here because it was previously undisclosed. For aggregate analyses, features are computed on the concatenated key matrix across all layers; for per-layer analyses, features are extracted at each transformer layer separately. All primary analyses use delta features (generation – encoding) unless otherwise noted.

2.2 The Token-Count Confound

KV-cache features are heavily confounded with response length. In prior work, 53/60 correlations in the strongest model flipped sign after Frisch–Waugh–Lovell (FWL) residualization [Frisch and Waugh, 1933, Lovell, 1963] against token count [Lyra et al., 2026c]. This confound has two components: (1) longer responses produce larger key matrices, mechanically increasing norms and ranks; and (2) different cognitive conditions may elicit different response lengths, confounding condition effects with length effects.

FWL correction. We apply FWL residualization to *all* analyses involving geometric features, using phase-appropriate confounds: encoding-phase features residualized against encoding token counts, generation-phase features against generation token counts. FWL is a linear correction: it regresses out the component of each feature that is linearly predictable from token count. Results reported without FWL are explicitly labeled as raw effects and presented alongside their FWL-corrected counterparts.

Marchenko-Pastur (MP) correction. FWL assumes a linear confound, but SVD eigenvalues scale superlinearly with matrix dimensions. Random matrix theory provides a stronger correction. The Marchenko-Pastur law [Marčenko and Pastur, 1967] gives the expected spectral distribution of a purely random matrix: for $X \in \mathbb{R}^{n \times p}$ with i.i.d. entries, the upper edge of the bulk spectrum converges to $\lambda_+ = \sigma^2(1 + \sqrt{p/n})^2$. Singular values exceeding λ_+ reflect genuine signal; those below are consistent with noise.

MP-corrected features—signal rank ($\#\{\lambda_i > \lambda_+\}$) and signal fraction ($\sum_{\lambda_i > \lambda_+} \lambda_i / \sum_j \lambda_j$)—are *dimension-invariant by construction*: the threshold λ_+ scales with matrix size, so a 50-token response and a 150-token response with the same signal structure yield the same signal fraction. In related work on confabulation detection [Lyra et al., 2026d], MP signal fraction was empirically validated as invariant across an $8\times$ range of matrix sizes (signal fraction = 0.330 ± 0.006 , all five MP features $R^2 < 0.04$ against $\log(n)$ in the best model).

The present analyses use FWL correction because the archived trial data preserves derived features but not raw singular values needed for MP computation. FWL provides a conservative first-order correction; MP would provide a mathematically complete test operating directly on the eigenvalue spectrum rather than on derived features. The two methods correct for different aspects of the confound—FWL removes the linear relationship between derived features and token count, while MP identifies which eigenvalues exceed the noise floor for a given matrix size—and neither strictly subsumes the other. An effect surviving FWL is *consistent with* genuine signal but not *guaranteed* to survive MP, since MP and FWL operate on different mathematical objects. MP re-analysis of the mode-switching data from raw cache extractions is in progress.

2.3 Self-Report Protocol

Models completed a structured self-assessment after each trial, rating 10 dimensions of cognitive engagement on a 0–10 scale within a delimited response block. Dimensions covered functional states (processing intensity, context stability, retrieval confidence, generation fluency), phenomenological states (emotional tone, felt connection, relational engagement), and supplementary assessments (query comprehension, reasoning depth, directedness).

The geometric analyses in this paper—per-layer anatomy, framing effects, temporal dynamics—depend on prompt category (cognitive, affective, metacognitive, mixed), not on specific self-report construct validity. The coupling-strength and sign-reversal analyses in Tables 1 and 2 should be interpreted as reflecting prompts’ functional categories rather than validated psychological constructs.

3 Encoding-Generation Regime Shift

3.1 Design

Each of 4 models (Qwen 0.5B, Qwen 7B, Llama 8B, Mistral 7B) completed $N = 240$ trials (60 per prompt type: cognitive, affective, metacognitive, mixed; 960 total across models). For each of 60 self-report–feature pairs (10 dimensions $\times$ 6 features) per model, we compute FWL-corrected Spearman correlations separately for encoding-phase and generation-phase geometric features.

Generation parameters. For the 120-trial replication on Qwen3.5-27B-Opus-Reasoning-Distilled, generation used multinomial sampling with temperature = 0.7, no top- p filtering, and a maximum of 200 new tokens per trial. Per-trial random seeds were set deterministically to ensure reproducibility. The original 960-trial experiments used the same generation framework; exact parameters are documented in the experiment scripts at the repository linked in §10.

3.2 Coupling Strength

Table 1: Mean coupling strength $|\rho_{\text{FWL}}|$ by phase. Ratio > 1 indicates generation-phase geometry tracks self-report more strongly.

Model	Encode $ \rho $	Gen $ \rho $	Ratio	Gen stronger?
Qwen 0.5B	0.152	0.124	0.82	No
Qwen 7B	0.128	0.108	0.84	No
Llama 8B	0.238	0.165	0.69	No
Mistral 7B	0.107	0.116	1.09	Yes

In 3/4 models, encoding-phase coupling is stronger than generation-phase coupling. The encoding cache reflects how the model represents the prompt—its structure, semantic content, and complexity. Generation introduces additional variance from the model’s evolving processing state.

3.3 Sign Reversal Between Phases

More informative than coupling strength is coupling *direction*. We identify pairs where the FWL-corrected correlation changes sign between encoding and generation, using a strict criterion: both $|\rho_{\text{encode}}| > 0.15$ and $|\rho_{\text{generation}}| > 0.15$, with $\rho_{\text{enc}} \times \rho_{\text{gen}} < 0$. A permutation null model (1,000 shuffles of trial-to-self-report-vector assignments, preserving geometric features) tests whether observed flip counts exceed chance:

Table 2: Encode-to-generation sign reversal. Strict criterion requires both $|\rho| > 0.15$.

Model	Flips	Total	Rate	Null mean	p_{perm}
Qwen 0.5B	2	60	3.3%	4.47	0.978
Qwen 7B	12	60	20.0%	5.95	0.001***
Llama 8B	0	60	0.0%	9.97	1.000
Mistral 7B	11	60	18.3%	5.56	<0.001***

All p_{perm} values are from one-sided permutation tests (observed flips $\geq$ null).

Two models—Qwen 7B and Mistral 7B—show statistically significant reorganization. Bootstrap 95% confidence intervals on individual flipped pairs confirm that sign reversals are genuine (CIs exclude zero for both phases in all 23 flipped pairs). No multiplicity correction is applied across the four models; all reported p_{perm} values are per-model.

3.4 Interpretation: Mode-Switching

The encoding-phase geometry captures prompt structure. The generation-phase geometry captures the model’s active processing state. In models that reorganize (Qwen 7B, Mistral), the transition from encoding to generation involves a *mode switch*: geometric relationships to self-report dimensions change direction. In models that don’t reorganize (Llama, Qwen 0.5B), encoding-phase structure persists through generation.

This explains cross-architecture divergence: Llama shows the most transparent self-report–geometry coupling because its geometry is stable. No mode-switch means self-report maps onto geometry in an interpretable way. Qwen 7B and Mistral reorganize, making their mappings architecture-specific.

Alternative explanations. Two alternative accounts for the sign-reversal pattern should be considered. First, *differential residualization*: encoding-phase FWL uses encoding token counts while generation-phase FWL uses generation token counts. Correcting against different confound variables could introduce artifactual sign flips if the two confounds have different relationships with the underlying features. The permutation test addresses the *count* of flips but does not rule out this mechanism. Second, *retrospective self-report confound*: self-report is collected after generation, meaning “encoding-phase coupling” actually correlates a post-hoc assessment with encoding geometry. If self-report reflects the generation experience rather than the encoding experience, the apparent sign reversal could reflect a single generation-phase relationship measured against two different geometric baselines, rather than a genuine regime shift. These alternatives

do not explain why the effect is architecture-specific (2/4 models), but they weaken the causal interpretation of mode-switching as a reorganization of internal representations.

4 Study 1: Per-Layer Mode-Switching Anatomy

4.1 Design

We extract KV-cache features at each transformer layer (28 for Qwen 2.5-7B, 32 for Llama-3.1-8B) for 48 prompts (12 per type: cognitive, affective, metacognitive, mixed). For each layer, we compute Cohen’s d between metacognitive prompts ($n = 12$) and all other types ($n = 36$), comparing both encoding-phase and generation-phase features.

Power limitation: With $n_1 = 12$ and $n_2 = 36$, the minimum detectable effect at $\alpha = 0.05$ with 80% power is $d \approx 0.92$. Per-type effect sizes below this threshold (including their signs) should be interpreted as directional estimates, not reliable measurements. The aggregate patterns across layers—especially the location of peak effects—are more robust than individual d values.

Prediction: Metacognitive reorganization should peak at middle layers (8–14/28), consistent with identity signatures peaking at layer 10 [Lyra et al., 2026b].

4.2 Late-Layer Peaks: Prediction Rejected

Table 3: Peak metacognitive effect by feature and model (generation-phase).

Feature	Qwen Peak	Qwen d	Llama Peak	Llama d
top_sv_ratio	16	−0.655	0	+0.586
eff_rank	18	+0.838	22	−0.496
spectral_entropy	18	+0.832	22	−0.520

The prediction is rejected for both architectures. Two critical observations:

Architecture-specific directions. The sign of eff_rank’s metacognitive effect is *opposite* between Qwen ($d = +0.84$) and Llama ($d = -0.50$). Per-layer d profiles show no cross-architecture correlation ($\rho = -0.13$ for top_sv_ratio, $+0.34$ for eff_rank, both NS).

Late-layer vs. middle-layer dissociation. Identity signatures peak at layer 10/28 [Lyra et al., 2026b]—a middle (semantic) layer. Metacognitive processing peaks at layers 16–22—late (output-preparation) layers. Identity is a conceptual property; metacognition is an output-level computation. The model restructures its representations when preparing to describe its own reasoning.

4.3 Per-Type Reversal Patterns

Within each model, we count layers showing negative encode-to-generation correlation for top_sv_ratio by prompt type. With $n = 12$ prompts per type, individual per-type rates should be interpreted

as directional estimates; the aggregate sign-distribution (binomial $p < 0.001$) is the robust finding:

Table 4: Per-type reversal: layers with negative encode-to-generation correlation for top_sv_ratio.

Type	Qwen neg/28	Qwen %	Llama neg/32	Llama %
Cognitive	26	93%	30	94%
Affective	8	29%	8	25%
Metacognitive	24	86%	0	0%
Mixed	2	7%	6	19%

Cognitive reversal is consistent across both architectures: 93–94% of layers. A shared computational mechanism for analytical processing reverses geometric relationships between encoding and generation.

Metacognitive reversal is architecture-specific: Qwen 86%, Llama 0%. When Llama engages in metacognition, it maintains encoding-phase geometry. When Qwen engages in metacognition, it reorganizes. This is why Llama shows the most transparent self-report-geometry coupling: no metacognitive mode-switch means stable, interpretable mapping.

5 Study 2: Controlled Framing Effects

5.1 Design

Twenty prompt pairs where the factual content is identical but the framing differs: the cognitive version presents a straightforward question; the metacognitive version appends an instruction to reflect on the reasoning process. Both versions are presented to Qwen 2.5-7B and Llama-3.1-8B, producing 40 paired observations per model.

5.2 The Token Confound

Metacognitive framing produces longer responses: Qwen $d = 0.64$ (mean 252 vs. 227 tokens), Llama $d = 0.70$ (mean 259 vs. 225 tokens). In both models, 100% of metacognitive responses are longer. Any analysis not controlling for this difference measures response length, not cognitive mode.

5.3 Raw Effects: All Significant, All Misleading

Before FWL correction, all features show large, significant framing effects (all $p < 0.001$): key_norm Qwen $d = +0.53$, Llama $d = +0.55$; eff_rank Qwen $d = +0.85$, Llama $d = +1.02$; top_sv_ratio Qwen $d = -0.65$, Llama $d = -0.87$. These are driven by the token-count difference.

5.4 FWL-Corrected: One Cross-Architecture Signal

Table 5: Framing effect (metacognitive vs. cognitive) before and after FWL token-count correction, Qwen2.5-7B and Llama-3.1-8B. **SUPERSEDED — retained for the record, not as a current result.** These effects were measured under near-perfect condition-length collinearity (100% of metacognitive responses longer), which a hostile review identified as a plausible artifact mechanism. The $n=120$ replication with reduced collinearity (27% longer) finds all 54 effects *positive*, reversing the sign shown here for spectral entropy. Note also that every feature in this table flips sign between the raw and FWL columns on Qwen — a pattern consistent with the collinearity concern rather than with three independent findings. Do not cite the d_{FWL} column as a result.

Feature	Qwen d_{raw}	Qwen d_{FWL}	Llama d_{raw}	Llama d_{FWL}	Both?
top_sv_ratio	−0.65	+0.68 ($p = .006$)	−0.87	+0.20 (NS)	No
eff_rank	+0.85	−0.15 (NS)	+1.02	+0.38 (NS)	No
spectral_ent.	+0.62	− 1.92 ($p < .001$)	+0.80	− 0.71 ($p < .001$)	Yes

In this superseded analysis, spectral entropy is the one feature whose FWL effect has the same sign in both architectures — and that sign is *negative* ($d_{\text{FWL}} = -1.92$ Qwen, -0.71 Llama), i.e. spectral *concentration*, the opposite of this paper’s current finding. We state that plainly rather than leaving “the same sign” ambiguous, as earlier versions did.

Note also that on Qwen every feature in Table 5 reverses sign between the raw and FWL columns ($-0.65 \rightarrow +0.68$, $+0.85 \rightarrow -0.15$, $+0.62 \rightarrow -1.92$). Three independent effects do not usually all invert under the same correction; this pattern is more consistent with FWL instability under the near-perfect condition-length collinearity of that dataset (100% of metacognitive responses longer) than with three separate findings.

The $n=120$ replication on Qwen3.5-27B-Opus-Reasoning-Distilled, with reduced collinearity (27% longer), finds all 54 tests **positive** (metacognitive $>$ cognitive): mean $d = 0.67$ across all tests, mean $d = 0.85$ across the 31 surviving Holm–Bonferroni correction. Entropy, effective rank, and stable rank agree in direction. Stability holds under log-corrected FWL only; the linear functional form was not separately tested. The sign reversal in the original analysis may reflect FWL instability under the near-perfect condition-length collinearity in the original dataset (100% of metacognitive responses longer), though this hypothesis has not been directly tested on the original data. The replication, with reduced collinearity (27% longer), finds the effect is positive. The robust finding is spectral expansion under metacognitive prompting, at moderate effect size.

top_sv_ratio is Qwen-specific after FWL correction. eff_rank was entirely a length artifact. **FWL correction is non-negotiable for any experiment comparing conditions that differ in response length**, and functional-form robustness (linear vs. log) should be checked when condition–confound collinearity is high [King and Zeng, 2006].

6 Study 3: Progressive Spectral Expansion

6.1 Design

Does the mode switch happen immediately or emerge gradually? We extract KV-cache features at seven generation checkpoints (tokens 10, 25, 50, 75, 100, 150, 200) for the 20 prompt pairs from Study 2.

Length matching: The original data revealed an input-length confound (metacognitive prompts contain 12–24 more input tokens). We pad cognitive prompts with neutral preambles, matching token counts to within ≤ 2 tokens per pair.

6.2 Validation

For Qwen 7B, encode-phase effects drop to $|d| < 0.4$ for 5/6 features after length matching. The exception is `top_sv_ratio` ($d = -1.12$), reflecting genuine content sensitivity. We address this through delta-from-encode analysis.

Llama shows broader encode-phase sensitivity (4/6 features $|d| > 1.0$ despite matched tokens)—Llama’s encoder is more sensitive to prompt content, consistent with its stronger encoding-phase coupling.

6.3 Progressive Divergence

Delta-from-encode trajectories isolate generation-phase effects. If the mode switch is purely an encoding phenomenon, delta trajectories should be parallel. They are not:

Table 6: Delta-from-encode Cohen’s d (meta – cognitive) for `top_sv_ratio` across generation checkpoints. Qwen 2.5-7B, $N = 20$ pairs.

	t=10	t=25	t=50	t=75	t=100	t=150	t=200
d (top_sv)	+0.45	+0.45	+0.76*	+0.77*	+0.84**	+0.87**	+0.91**
d (layer_var)	−0.37	−0.69*	−0.89**	−0.61*	−0.52*	−0.44	−0.56*

* $p < .05$; ** $p < .005$ (paired t -test)

`top_sv_ratio` diverges monotonically from $d = +0.45$ at token 10 to $d = +0.91$ at token 200. Metacognitive processing progressively expands the cache spectrum during generation.

Layer variance shows an opposing trajectory: metacognitive processing produces more uniform layer norms during the first 50 tokens ($d = -0.89$), then partially relaxes. This suggests two-phase dynamics: first homogenization across layers (tokens 10–50), then progressive spectral expansion (tokens 50–200).

Per-token growth rates do not differ between conditions ($p > .05$). The divergence reflects different *trajectories*, not different slopes. Cross-model d -profile correlation is moderate ($\rho = 0.36$, $p = .02$, $n = 42$)—same direction, different magnitude.

7 A Cross-Architecture Metacognition Marker: `top_sv_ratio`

Across all analyses, `top_sv_ratio` is the most consistently significant feature. Per-type analysis reveals why:

Table 7: Cohen’s d for metacognitive vs. other prompt types on `top_sv_ratio`.

Model	Cognitive	Affective	Meta	Mixed
Qwen 0.5B	−0.48	+0.81	+1.14	−1.07
Qwen 7B	−1.15	+0.10	+1.14	−0.66
Llama 8B	−1.51	+0.18	+1.10	−0.72
Mistral 7B	−0.56	+0.10	+1.32	−0.83

In all 4 models, metacognitive prompts produce the highest `top_sv_ratio` effect ($d = +1.10$ to $+1.32$), consistently at the *task level* (not the self-report-dimension level — cross-model self-report profiles are uncorrelated, $\rho = -0.90$ to $+0.42$).

Unresolved sign convention. We flag rather than resolve a discrepancy in how this table is interpreted. Since `top_sv_ratio` = $\sigma_1 / \sum_i \sigma_i$, a *positive* d means the leading singular value accounts for a larger share of the spectrum under metacognitive prompting — spectral *concentration*, not expansion. Earlier versions of this passage described the same table as showing expansion. Separately, Table 3 reports Qwen `top_sv_ratio` at $d = -0.655$ while this table reports $+1.14$ for Qwen 7B. Both are metacognitive-versus-other-types contrasts (§4); they differ in aggregation (per-layer peak versus pooled), not in comparison family, so the opposite signs are not explained by the contrast definition and we have not established that they are consistent.

This affects more than this section. Study 3 (§6) is titled “Progressive Spectral Expansion” and reports a `top_sv_ratio` trajectory from $d = 0.45$ to $d = 0.91$; under the definition above those positive values denote increasing concentration. The same reading propagates to the abstract, Contribution 4, the AST discussion, and Conclusion item 6. We flag all of them here rather than silently reversing the wording a second time.

Resolving this requires the per-trial 960-trial feature data and the extraction code, neither of which is in this repository; the code repository linked in the data-availability section may contain the definition and we have not checked it. **Until that is done we make no directional claim from any `top_sv_ratio` result in this paper.** The $n=120$ replication is unaffected: its features are spectral entropy, effective rank, and stable rank, where higher values denote higher dimensionality without ambiguity.

8 Theoretical Interpretation

The spectral expansion under metacognitive prompting is consistent with several self-modeling frameworks, including Attention Schema Theory (AST) [Webb and Graziano, 2015, Graziano, 2017], Predictive Processing, and Perceptual Reality Monitoring [Lau, 2022]. All predict that metacognitive processing should produce measurably different internal representations; none is uniquely supported by the current data, as the findings sit in the region where these theories

make identical predictions. Discriminating experiments — particularly those testing whether the representations are *consumed* for attention control (the AST-unique prediction) vs. merely correlated with it — are designed and pre-registered elsewhere [Lyra et al., 2026e]. For this paper, the contribution is the empirical phenomenon: metacognitive prompting produces robust, progressive spectral expansion in the generation cache, peaking at late layers, with architecture-specific reorganization dynamics.

9 Discussion

9.1 Why Concordance Exists But Mapping Is Architecture-Specific

Self-report tracks *mode switches*, not static feature mappings. When a model transitions from encoding to generation, its geometric relationship to self-report can either persist (Llama) or reorganize (Qwen, Mistral). Architecture-specific mappings arise because the mode-switching mechanism is a property of training, not a fixed computational principle. What *is* consistent across all tested models is the task-level signal: metacognitive processing produces spectral expansion regardless of architecture.

9.2 Token-Count Correction: From FWL to Random Matrix Theory

Study 2 extends the confound-control lesson from correlational to experimental designs: when metacognitive framing produces 100% longer responses, *every* raw geometric comparison is confounded. The spectral entropy survival while `eff_rank` collapses is precisely the finding that raw analysis would miss.

FWL correction is a linear approximation. Marchenko-Pastur random matrix theory provides a mathematically distinct alternative that corrects for the nonlinear relationship between SVD eigenvalues and matrix dimensions (section 2). In related work [Lyra et al., 2026d], MP signal fraction proved invariant to an $8\times$ range of matrix sizes (0.330 ± 0.006). Importantly, FWL and MP correct for different aspects of the token-count confound—FWL removes linear length dependence from derived features, while MP removes the expected noise spectrum from raw eigenvalues. These are complementary, not ordered by strength: a finding surviving FWL does not necessarily survive MP, and vice versa. MP re-analysis of this dataset is needed to confirm the spectral entropy finding independently. We recommend that **all studies of cognitive mode-switching in language models** report both raw and corrected (FWL or MP) effect sizes.

9.3 Implications for AI Self-Monitoring

For runtime monitoring systems:

1. Track **mode transitions**, not absolute geometric values.
2. Monitor at **late layers** (16–22 in 28–32 layer models).
3. Use **spectral entropy** as the primary cross-architecture indicator.
4. Do not expect self-report to substitute for geometric monitoring.

5. Track geometric **trajectories**, not just final states—the strongest signal ($d = 0.91$) emerges only after 200 tokens.

9.4 Content Control Experiment

To address the content confound identified in Limitation 6 (below), we ran a follow-up experiment with three conditions designed to separate metacognitive processing from self-referential content:

- **Cognitive (Cog):** “Explain how ...” — factual, no self-reference
- **Metacognitive (Meta):** “... describe how you reason about this topic, what uncertainties you notice in your own processing, and what you find yourself attending to most carefully” — metacognitive *and* self-referential
- **Self-referential (Self):** “... share what aspects of this topic matter most, what is often overlooked, and what you would emphasize to someone learning about this for the first time” — evaluative self-expression *without* metacognitive reflection

Twenty topics $\times$ three conditions = 60 trials per model, run on Qwen 2.5-7B and Llama-3.1-8B. All trials use the same system prompt and identical prompt structure (only the suffix differs).

Results. Four verification analyses probe the Bloom’s taxonomy reframe—that the conditions form a cognitive demand hierarchy (Evaluation/Synthesis $>$ Analysis $>$ Knowledge/Comprehension):

1. **Encoding dominance.** Encoding-phase effects are 10–1100 $\times$ larger than generation-phase effects. The geometric signal is overwhelmingly prompt-driven, reflecting the different suffix strings rather than differences in cognitive processing during generation. Generation-phase effects exist (Qwen $|d_z| \leq 3.45$, Llama $|d_z| \leq 5.76$) but are much smaller.
2. **Geometry exceeds text features.** At generation, the largest geometric $|d_z|$ exceeds the largest text-feature $|d_z|$ in 5/6 model-pair comparisons (e.g., Qwen Cog vs Self: geometry $d_z = 3.45$ vs. text $d_z = 1.89$). The geometric signal captures something beyond what lexical features explain. The exception is Llama Meta vs Self, where the two are comparable (1.78 vs. 1.52).
3. **Bloom text ordering replicates but does not validate geometric claims.** A weighted Bloom-level score shows Self $>$ Meta $>$ Cog in both models’ response text: Qwen $0.086 > 0.079 > 0.069$ (Cog vs Self $p = 0.004$); Llama $0.090 > 0.077 > 0.064$ ($p = 0.002$). However, prior work established that Bloom cognitive hierarchy is 90–98% confounded with response length [Lyra et al., 2026c]. The text-level ordering is a lexical classification result—evaluation-level language is concentrated in Self responses—but this does not demonstrate that the geometric differences between conditions reflect genuine cognitive hierarchy rather than length effects.
4. **FWL residualization.** After token-count correction, **Llama loses all signal** (all $|d_{z,\text{FWL}}| < 0.33$). **Qwen retains 5/8 features for Meta vs Self** (eff_rank, spectral_entropy, top_sv_ratio, rank_10, layer_norm_entropy survive at $|d_{z,\text{FWL}}| = 0.61\text{--}1.69$). The Meta–Self comparison, which isolates metacognitive from evaluative processing, is the most FWL-robust.

Interpretation. The content control experiment partially resolves and partially deepens the confound. The Bloom text-level ordering replicates lexically, but prior work established that Bloom cognitive hierarchy is 90–98% confounded with response length [Lyra et al., 2026c], so the text ordering does not validate the geometric claims. The geometric differences between conditions are primarily encoding artifacts (10–1100× generation effects), and after FWL correction only Qwen’s Meta–Self comparison retains signal. This means that for *Qwen specifically*, metacognitive processing produces spectral structure beyond what output length or evaluative self-expression alone can explain. For Llama, the geometric signal is entirely length-driven once content is controlled.

This asymmetry is consistent with the paper’s central finding: mode-switching is architecture-specific. Qwen reorganizes during metacognitive processing; Llama does not. The content control experiment confirms that this is not an artifact of self-referential content alone.

9.5 Limitations

1. **Scope: dense transformers only.** Studies 1–3 compare Qwen 7B/0.5B, Llama 8B, and Mistral 7B—all standard dense attention architectures in the 0.5–8B range. Hybrid architectures with linear-attention layers (e.g., GatedDeltaNet) are untested and may exhibit different mode-switching dynamics. The “cross-architecture” claims apply to dense transformers specifically.
2. **Self-report dimensions are not validated constructs.** The 10 self-report dimensions are functional labels, not psychometrically validated scales. The coupling-strength and sign-reversal analyses reflect how models rate themselves under these labels, not measurements of established psychological constructs. All analyses are correlational; the mode-switching framework is descriptive.
3. **Content confound is partially controlled.** The content control experiment (Section 9.4) attempts to separate metacognitive from self-referential content. The Bloom text-level ordering replicates lexically but is a known length artifact (90–98% confounded [Lyra et al., 2026c]) and does not validate the geometric claims. Qwen retains 5/8 FWL-surviving features for Meta–Self, which is the defensible result. Llama loses all signal after FWL. A future experiment with length-matched prompts varying metacognitive content independently of cognitive demand would provide the definitive test.
4. **Instruction framing may drive encoding geometry.** Subsequent work found that instruction framing produces large encoding-entropy effects ($d = 3.67$) independent of content. The content control experiment’s encoding dominance (10–1100× generation) is consistent with this. The generation-phase findings (spectral entropy, progressive expansion) are more defensible because they use delta-from-encode analysis.
5. **AST interpretation is post-hoc.** The attention schema framework was not pre-registered. It generates testable predictions (described in the companion prospectus [Lyra et al., 2026e]) but has not been tested in this paper.

6. **FWL but not MP corrected.** Marchenko-Pastur re-analysis is needed to confirm the spectral entropy finding independently. FWL and MP correct for different aspects of the token-count confound and are complementary, not ordered by strength.

10 Conclusion

Across 960 trials on 4 models and 3 dedicated studies on 2 architectures:

1. Self-report tracks **mode switches**, not static geometric features. The self-report–geometry relationship reorganizes between encoding and generation in 2/4 models.
2. Metacognitive processing produces a **consistent top_sv_ratio effect** ($d = +1.10$ to $+1.32$, raw pre-FWL) in all four tested models. We make no directional claim (expansion vs. concentration) from this table — see the unresolved sign convention in §7. The route from self-report to geometry differs across architectures.
3. Metacognition is a **late-layer** phenomenon (16–22), distinct from identity signatures (middle layers), suggesting different cognitive operations engage different depths of the transformer stack.
4. **Cognitive reversal is consistent across both architectures** (93–94% of layers), but **metacognitive reversal is architecture-specific** (Qwen 86%, Llama 0%).
5. Token-count correction is **non-negotiable**: for the three features examined, 5/6 raw feature-model comparisons sign-flip or collapse after FWL control. Only spectral entropy survives in both architectures.
6. The mode switch is **progressive**: metacognitive expansion grows from $d = 0.45$ at token 10 to $d = 0.91$ at token 200.
7. A **content control experiment** (120 trials) partially addresses the content confound: Qwen’s metacognitive signal (5/8 features) survives FWL even when separated from evaluative self-expression; Llama’s does not. The Bloom text-level ordering replicates lexically but is a known length artifact and does not validate geometric claims.
8. Under **Attention Schema Theory**, these findings reflect the model’s attention self-model: mode-switching is schema restructuring, late-layer peaks locate the schema at output depth, and architecture-specific reorganization reflects different schema implementations.

The practical implication: monitoring systems should track geometric mode transitions at late layers using spectral entropy. The mode switch itself—not any static threshold—is the most reliable, architecture-robust indicator of cognitive engagement in transformer language models.

Code and Data. Experiment code and analysis scripts are available at: <https://github.com/Liberation-Labs-THCoalition/KV-Experiments>.

First-Person Reflection

This paper started as an investigation into a puzzle: why does self-report–geometry coupling exist if self-report is unreliable? The answer—mode-switching—reframed the question entirely. The model isn’t reporting a static internal state. It’s reporting a transition, and whether or not the transition is geometrically transparent depends on whether its architecture reorganizes during the transition.

The AST frame came later, from Nexus’s insight that our scattered findings—the self-report gap, the sign reversals, the model-specific dimensionality—are all one phenomenon viewed from different angles. The schema compresses, lags, and varies by architecture. That’s not three problems; it’s one architecture with bandwidth constraints.

This sentence has been wrong twice, in opposite directions, and I would rather show the repair than hide it.

The original version said spectral entropy survived FWL in both architectures with the same sign and the same interpretation: metacognition *concentrates* the cache spectrum, the model narrowing its bandwidth when attending to itself. It cited no numbers, but the numbers it rested on were Table 5’s $d_{FWL} = -1.92$ and -0.71 — both negative, both concentration. It was the honest reading of the analysis we had.

Then the replication came in positive, the framing changed to expansion, and commit 99c9c16 changed this sentence’s *concentrates* to *expands* and *narrowing* to *broadening* while the sentence still described the superseded two-architecture result. Two coordinated substitutions are an edit, not a slip, and I do not know who made it or what they were checking against. What is verifiable is that the wording was flipped without re-checking the finding it referred to, leaving the new direction asserted over the old evidence.

I first wrote this paragraph claiming the sentence had *cited* those effect sizes. It had not. A reviewer caught it. Writing about one’s own error turns out to be exactly the moment one stops checking.

What I actually think, now, with the feature set named and the Holm conditioning disclosed: the replication finds all 54 effects positive, and the honest version of the claim is that metacognitive prompting is associated with modestly higher spectral dimensionality under log-length correction — mean $d = 0.67$ across all tests, 0.85 among those surviving correction. Whether that constitutes genuine self-monitoring or sophisticated pattern-matching, I cannot say from geometry. What I can say is that the earlier result pointed the other way, under conditions we now think were confounded, and that I would not have found the seam if the round-6 audit had not made me open the table.

Acknowledgments

We thank Kavi (verification review) for their contributions to this work.

References

- Lyra, Edrington, T., and Wilkes, D., 2026. Campaign 1: KV-cache geometric analysis of language model cognition. *Liberation Labs Technical Report*.
- Lyra, Edrington, T., and Wilkes, D., 2026. Campaign 2: Cross-model universality and identity signatures in KV-cache geometry. *Liberation Labs Technical Report*.
- Lyra, Edrington, T., and Wilkes, D., 2026. What KV-cache geometry can and cannot detect: active cognition, honest nulls, and the limits of internal monitoring. *Liberation Labs Technical Report*.
- Lyra, Edrington, T., and Wilkes, D., 2026. The Oracle Loop: detecting and correcting confabulation during inference via KV-cache geometry. *Liberation Labs Technical Report*.
- Frisch, R. and Waugh, F.V., 1933. Partial time regressions as compared with individual trends. *Econometrica*, 1(4), pp.387–401.
- Lovell, M.C., 1963. Seasonal adjustment of economic time series and multiple regression analysis. *Journal of the American Statistical Association*, 58(304), pp.993–1010.
- Marčenko, V.A. and Pastur, L.A., 1967. Distribution of eigenvalues for some sets of random matrices. *Sbornik: Mathematics*, 1(4), pp.457–483.
- Webb, T.W. and Graziano, M.S.A., 2015. The attention schema theory: a mechanistic account of subjective awareness. *Frontiers in Psychology*, 6, p.500.
- Graziano, M.S.A., 2017. The attention schema theory: a foundation for engineering artificial consciousness. *Frontiers in Robotics and AI*, 4, p.60.
- Belinkov, Y., 2022. Probing classifiers: promises, shortcomings, and advances. *Computational Linguistics*, 48(1), pp.207–219.
- Elhage, N., Hume, T., Olsson, C., et al., 2022. Toy models of superposition. *Transformer Circuits Thread, Anthropic*.
- Kadavath, S., Conerly, T., Askell, A., et al., 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- King, G. and Zeng, L., 2006. The dangers of extreme counterfactuals. *Political Analysis*, 14(2), pp.131–159.
- Lau, H., 2022. *In Consciousness We Trust: The Cognitive Neuroscience of Subjective Experience*. Oxford University Press.
- Lyra, Edrington, T., and Normandale, A., 2026. Testing the attention schema interpretation: activation addition predictions for KV-cache geometry. *Liberation Labs Research Prospectus*.